

House Price Linear Regression model in R Studio

GAURAV SIWAL (M.TECH)

2020-05-28

```
df = read.csv("C:/Users/CEA/Desktop/ML by R Alison/Complete ML in R/1. Linear Regression/House_Price.csv", header = TRUE)
```

```
View(df)
```

```
summary(df)
```

```
##      price      crime_rate      resid_area      air_qual
## Min.   : 5.00   Min.   : 0.00632   Min.   :30.46   Min.   :0.3850
## 1st Qu.:17.02   1st Qu.: 0.08204   1st Qu.:35.19   1st Qu.:0.4490
## Median :21.20   Median : 0.25651   Median :39.69   Median :0.5380
## Mean   :22.53   Mean   : 3.61352   Mean   :41.14   Mean   :0.5547
## 3rd Qu.:25.00   3rd Qu.: 3.67708   3rd Qu.:48.10   3rd Qu.:0.6240
## Max.   :50.00   Max.   :88.97620   Max.   :57.74   Max.   :0.8710
##
##      room_num      age      dist1      dist2
## Min.   :3.561   Min.   : 2.90   Min.   : 1.130   Min.   : 0.920
## 1st Qu.:5.886   1st Qu.: 45.02   1st Qu.: 2.270   1st Qu.: 1.940
## Median :6.208   Median : 77.50   Median : 3.385   Median : 3.010
## Mean   :6.285   Mean   : 68.57   Mean   : 3.972   Mean   : 3.629
## 3rd Qu.:6.623   3rd Qu.: 94.08   3rd Qu.: 5.367   3rd Qu.: 4.992
## Max.   :8.780   Max.   :100.00   Max.   :12.320   Max.   :11.930
##
##      dist3      dist4      teachers      poor_prop      airport
## Min.   : 1.150   Min.   : 0.730   Min.   :18.00   Min.   : 1.73   NO :227
## 1st Qu.: 2.232   1st Qu.: 1.940   1st Qu.:19.80   1st Qu.: 6.95   YES:279
## Median : 3.375   Median : 3.070   Median :20.95   Median :11.36
## Mean   : 3.961   Mean   : 3.619   Mean   :21.54   Mean   :12.65
## 3rd Qu.: 5.407   3rd Qu.: 4.985   3rd Qu.:22.60   3rd Qu.:16.95
## Max.   :12.320   Max.   :11.940   Max.   :27.40   Max.   :37.97
##
##      n_hos_beds      n_hot_rooms      waterbody      rainfall
## Min.   : 5.268   Min.   : 10.06   Lake           : 97   Min.   : 3.00
## 1st Qu.: 6.635   1st Qu.: 11.19   Lake and River: 71   1st Qu.:28.00
## Median : 7.999   Median : 12.72   None           :155   Median :39.00
## Mean   : 7.900   Mean   : 13.04   River          :183   Mean   :39.18
## 3rd Qu.: 9.088   3rd Qu.: 14.17           3rd Qu.:50.00
## Max.   :10.876   Max.   :101.12           Max.   :60.00
## NA's      :8
## bus_ter      parks
## YES:506   Min.   :0.03329
##           1st Qu.:0.04646
##           Median :0.05351
##           Mean   :0.05445
```

```
##          3rd Qu.:0.06140
##          Max.    :0.08671
##

# treatment for missing data
mean_val = mean(df$n_hos_beds, na.rm = TRUE)
df$n_hos_beds[is.na(df$n_hos_beds)] = mean_val

summary(df$n_hos_beds)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  5.268   6.659   7.963   7.900   9.076  10.876

# outlier treatment
uv = 3* quantile(df$crime_rate, 0.99)
df$crime_rate[df$crime_rate > uv]= uv
summary(df$crime_rate)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.00632  0.08204  0.25651  3.61352  3.67708 88.97620

lv = 0.3 * quantile(df$rainfall, 0.01)
df$rainfall[df$rainfall < lv]= lv
summary(df$rainfall)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    6.00   28.00   39.00   39.19   50.00   60.00

head(df)

##   price crime_rate resid_area air_qual room_num  age dist1 dist2 dist3
## 1  24.0    0.00632    32.31    0.538    6.575 65.2  4.35  3.81  4.18
## 2  21.6    0.02731    37.07    0.469    6.421 78.9  4.99  4.70  5.12
## 3  34.7    0.02729    37.07    0.469    7.185 61.1  5.03  4.86  5.01
## 4  33.4    0.03237    32.18    0.458    6.998 45.8  6.21  5.93  6.16
## 5  36.2    0.06905    32.18    0.458    7.147 54.2  6.16  5.86  6.37
## 6  28.7    0.02985    32.18    0.458    6.430 58.7  6.22  5.80  6.23
## teachers poor_prop airport n_hos_beds n_hot_rooms waterbody rainfall
## bus_ter
## 1    24.7     4.98     YES     5.480     11.1920     River      23
## 2    22.2     9.14     NO      7.332     12.1728     Lake      42
## 3    22.2     4.03     NO      7.394    101.1200     None      38
```

```

## 4      21.3      2.94      YES      9.268      11.2672      Lake      45
YES
## 5      21.3      5.33      NO      8.824      11.2896      Lake      55
YES
## 6      21.3      5.21      YES      7.174      14.2296      None      53
YES
##           parks
## 1 0.04934731
## 2 0.04614563
## 3 0.04576397
## 4 0.04715060
## 5 0.03947400
## 6 0.04590965

# average dist
df$avg_dist = (df$dist1 + df$dist2+ df$dist3+ df$dist4)/4
head(df)

##   price crime_rate resid_area air_qual room_num  age dist1 dist2 dist3
dist4
## 1  24.0    0.00632    32.31    0.538    6.575 65.2  4.35  3.81  4.18
4.01
## 2  21.6    0.02731    37.07    0.469    6.421 78.9  4.99  4.70  5.12
5.06
## 3  34.7    0.02729    37.07    0.469    7.185 61.1  5.03  4.86  5.01
4.97
## 4  33.4    0.03237    32.18    0.458    6.998 45.8  6.21  5.93  6.16
5.96
## 5  36.2    0.06905    32.18    0.458    7.147 54.2  6.16  5.86  6.37
5.86
## 6  28.7    0.02985    32.18    0.458    6.430 58.7  6.22  5.80  6.23
5.99
##   teachers poor_prop airport n_hos_beds n_hot_rooms waterbody rainfall
bus_ter
## 1    24.7    4.98      YES    5.480    11.1920    River      23
YES
## 2    22.2    9.14      NO    7.332    12.1728    Lake      42
YES
## 3    22.2    4.03      NO    7.394    101.1200    None      38
YES
## 4    21.3    2.94      YES    9.268    11.2672    Lake      45
YES
## 5    21.3    5.33      NO    8.824    11.2896    Lake      55
YES
## 6    21.3    5.21      YES    7.174    14.2296    None      53
YES
##           parks avg_dist
## 1 0.04934731  4.0875
## 2 0.04614563  4.9675
## 3 0.04576397  4.9675

```

```
## 4 0.04715060    6.0650
## 5 0.03947400    6.0625
## 6 0.04590965    6.0600

df= df[,c(-7,-8,-9,-10)]

# categorical value treatment
library('dummies')

## dummies-1.5.6 provided by Decision Patterns

df = dummy.data.frame(df)

## Warning in model.matrix.default(~x - 1, model.frame(~x - 1), contrasts =
FALSE):
## non-list contrasts argument ignored

## Warning in model.matrix.default(~x - 1, model.frame(~x - 1), contrasts =
FALSE):
## non-list contrasts argument ignored

df = df[,c(-9,-15)]

# train test split
library('caTools')

## Warning: package 'caTools' was built under R version 3.6.3

set.seed(0)
split = sample.split(df, SplitRatio = 0.8)
train = subset(df, split == 'TRUE')
test = subset(df, split == 'FALSE')

# linear regression
lreg = lm(price~.,data = train)
summary(lreg)

##
## Call:
## lm(formula = price ~ ., data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -10.113   -2.963   -0.800    1.964   26.580
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -4.113657    5.912565  -0.696   0.4870
## crime_rate   -0.067481    0.038063  -1.773   0.0771 .
## resid_area   -0.037266    0.067811  -0.550   0.5830
## air_qual    -14.791888    6.520164  -2.269   0.0239 *
```

```

## room_num          4.036105    0.484752    8.326 1.65e-15 ***
## age               -0.005714    0.015766   -0.362  0.7172
## teachers          0.888358    0.137164    6.477 3.00e-10 ***
## poor_prop        -0.652535    0.059251  -11.013 < 2e-16 ***
## airportYES        1.141855    0.531672    2.148  0.0324 *
## n_hos_beds         0.370085    0.178322    2.075  0.0386 *
## n_hot_rooms        0.029834    0.043419    0.687  0.4924
## waterbodyLake      0.258283    0.787181    0.328  0.7430
## `waterbodyLake and River` -0.709849    0.792252   -0.896  0.3708
## waterbodyRiver     -0.629699    0.629548   -1.000  0.3179
## rainfall           0.011073    0.021009    0.527  0.5985
## parks             31.387873   59.037418    0.532  0.5953
## avg_dist          -1.298917    0.216917   -5.988 5.04e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.998 on 369 degrees of freedom
## Multiple R-squared:  0.7287, Adjusted R-squared:  0.7169
## F-statistic: 61.95 on 16 and 369 DF,  p-value: < 2.2e-16

pred_train = predict(lreg, test)

mse = mean((pred_train - test$price)^2)
mse

## [1] 22.2732

# Other linear models for better prediction accuracy and model interpretability

# 1) Subset Selection Techniques

# 1.a) Best subset selection

library("leaps")

## Warning: package 'leaps' was built under R version 3.6.3

lm_sub_reg = regsubsets(price~., data = df, nvmax = 15)
summary(lm_sub_reg)$adjr2

## [1] 0.5479428 0.6410518 0.6792769 0.6909566 0.7088919 0.7120175 0.7144706
## [8] 0.7170322 0.7172295 0.7172748 0.7172916 0.7171582 0.7168656 0.7166040
## [15] 0.7162438

which.max(summary(lm_sub_reg)$adjr2)

## [1] 11

coef(lm_sub_reg, 11)

```

```
## (Intercept) crime_rate
air_qual
## -8.45195104 -0.07350473 -
21.46196860
## room_num teachers
poor_prop
## 4.15302687 0.98444383 -
0.55460143
## airportYES n_hos_beds `waterbodyLake and
River`
## 1.03063947 0.36294609 -
0.72313242
## rainfall parks
avg_dist
## 0.01790746 60.06200089 -
1.17184364
```

1.b) Forward stepwise selection

```
lm_fsub_reg = regsubsets(price~., data = df, nvmax = 15, method = 'forward')
summary(lm_fsub_reg)$adjr2
```

```
## [1] 0.5479428 0.6410518 0.6792769 0.6909566 0.7088919 0.7120175 0.7144706
## [8] 0.7170322 0.7172295 0.7172748 0.7172916 0.7171582 0.7168656 0.7166040
## [15] 0.7162438
```

```
which.max(summary(lm_fsub_reg)$adjr2)
```

```
## [1] 11
```

```
coef(lm_fsub_reg, 11)
```

```
## (Intercept) crime_rate
air_qual
## -8.45195104 -0.07350473 -
21.46196860
## room_num teachers
poor_prop
## 4.15302687 0.98444383 -
0.55460143
## airportYES n_hos_beds `waterbodyLake and
River`
## 1.03063947 0.36294609 -
0.72313242
## rainfall parks
avg_dist
## 0.01790746 60.06200089 -
1.17184364
```

```
# 1.c) Backward stepwise selection
```

```
lm_bsub_reg = regsubsets(price~., data = df, nvmax = 15, method = 'backward')
summary(lm_bsub_reg)$adjr2
```

```
## [1] 0.5479428 0.6410518 0.6792769 0.6909566 0.7088919 0.7120175 0.7144706
## [8] 0.7170322 0.7172295 0.7172748 0.7172916 0.7171582 0.7168656 0.7166040
## [15] 0.7162438
```

```
which.max(summary(lm_bsub_reg)$adjr2)
```

```
## [1] 11
```

```
coef(lm_bsub_reg, 11)
```

```
## (Intercept) crime_rate
air_qual -8.45195104 -0.07350473 -
21.46196860
## room_num teachers
poor_prop 4.15302687 0.98444383 -
0.55460143
## airportYES n_hos_beds `waterbodyLake and
River` 1.03063947 0.36294609 -
0.72313242
## rainfall parks
avg_dist 0.01790746 60.06200089 -
1.17184364
```

```
# 2) Shrinkage method ( Ridge and Lasso Technique)
```

```
# 2.1) ridge regression
```

```
library('glmnet')
```

```
## Warning: package 'glmnet' was built under R version 3.6.3
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.0
```

```
x= model.matrix(price~.,data = df)[,-1]
```


```
y = df$price
```

```
grid = 10^ seq(10,-2,length = 100)
```

```
lm_ridge = glmnet(x,y,alpha = 0, lambda = grid)
```

```
lm_ridgecv = cv.glmnet(x,y,alpha = 0, lambda = grid)
```

```
plot(lm_ridgecv)
```



```

optlambda = lm_ridgecv$lambda.min

pred_ridge = predict(lm_ridge, s= optlambda, newx = x)
mse_ridge = mean((pred_ridge - y)^2)
mse_ridge


## [1] 23.27537

TSS = mean((mean(y)- y)^2)
RSS = mean((pred_ridge - y)^2)
r2 = (1-(RSS/TSS))
r2

## [1] 0.7233922

# 2.2) Lasso regression
lm_lasso = glmnet(x,y, alpha = 1, lambda = grid)
lm_lassocv = cv.glmnet(x,y, alpha = 1, lambda = grid)
plot(lm_lassocv)

```

```

optlambda_lasso = lm_lassocv$lambda.min

pred_lasso = predict(lm_lasso, s= optlambda_lasso, newx = x)
mse_lasso = mean((pred_lasso - y)^2)
mse_lasso

## [1] 23.16464

TSS_lasso = mean((mean(y) - y)^2)
RSS_lasso = mean((pred_lasso - y)^2)
r2_lasso = 1 - (RSS_lasso / TSS_lasso)
r2_lasso

## [1] 0.7247081

```